
Sales Forecasting

Problem Statement:

Fresh Analytics aims to forecast the demand for various items across multiple restaurants. Accurate sales predictions enable better decision-making around staffing, inventory, and pricing strategies.

Datasets Used:

1. **restaurants.csv**: Details about restaurants (id, name)
2. **items.csv**: Details about items (id, name, kcal, cost, store_id)
3. **sales.csv**: Sales data (date, item, price, item_count)

Project Tasks:

1. Preliminary Analysis

- Load and inspect datasets.
- Handle missing values and outliers.
- Merge all datasets to a single DataFrame.

2. Exploratory Data Analysis (EDA)

- Analyze overall and time-based sales patterns.
- Weekday, monthly, and quarterly trends.
- Restaurant-wise and item-wise comparisons.
- High-performing restaurants and items.
- Caloric and cost analysis of top items.

3. Feature Engineering

- Extract temporal features: year, month, day, weekday, quarter
- Aggregate sales on different levels.

4. Model Building and Forecasting

- Models: Linear Regression, Random Forest, XGBoost.
- Train on historical data, test on the last 6 months.
- Evaluate using RMSE.

- Forecast for next year using the best model.

Insights:

- Identified top-selling items and restaurants.
- Analyzed temporal patterns in demand.
- Predicted future sales using ML algorithms.

Source Code:

1. Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor

from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
```

2. Load Datasets

```
df_rest = pd.read_csv("restaurants.csv")
df_items = pd.read_csv("items.csv")
df_sales = pd.read_csv("sales.csv")
```

3. Rename columns to avoid conflicts

```
df_items.rename(columns={'name': 'item_name'}, inplace=True)
```

```
df_rest.rename(columns={'name': 'store_name'}, inplace=True)
df_sales.rename(columns={'name': 'item_name'}, inplace=True)
```

```
# Optional: View column names
```

```
print("Sales Columns:", df_sales.columns)
print("Items Columns:", df_items.columns)
print("Restaurants Columns:", df_rest.columns)
```

```
# 4. Merge Datasets
```

```
df = df_sales.merge(df_items, on='item_name', how='left') \
    .merge(df_rest, left_on="store_id", right_on="id", suffixes=('_item', '_store'))
```

```
# 5. Convert 'date' and drop missing
```

```
df['date'] = pd.to_datetime(df['date'])
df.dropna(inplace=True)
```

```
# 6. Feature Engineering
```

```
df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
df['day'] = df['date'].dt.day
df['weekday'] = df['date'].dt.dayofweek
df['quarter'] = df['date'].dt.quarter
```

```
# 7. EDA
```

```
# Overall sales trend
```

```
plt.figure(figsize=(12,5))
df.groupby('date')['item_count'].sum().plot(title="Daily Sales Trend")
```

```
plt.xlabel("Date")
```

```
plt.ylabel("Total Items Sold")
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# Sales by weekday
```

```
plt.figure(figsize=(8,4))
```

```
sns.barplot(data=df, x='weekday', y='item_count', estimator=sum)
```

```
plt.title("Sales by Weekday")
```

```
plt.xlabel("Day of Week (0=Monday)")
```

```
plt.ylabel("Total Items Sold")
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# Top 10 items
```

```
top_items =
```

```
df.groupby('item_name')['item_count'].sum().sort_values(ascending=False).head(10)
```

```
print("Top 10 Items Sold:\n", top_items)
```

```
# Most profitable store
```

```
df['revenue'] = df['item_count'] * df['Price']
```

```
top_store = df.groupby('store_name')['revenue'].sum().sort_values(ascending=False)
```

```
print("Most Profitable Stores:\n", top_store)
```

```
# Most expensive item per store
```

```
expensive_items = df.groupby('store_name').apply(lambda x: x.sort_values('cost',  
ascending=False).iloc[0])
```

```
print("Most Expensive Item per Store:\n", expensive_items[['item_name', 'cost', 'kcal']])
```

8. Modeling and Forecasting

Aggregate daily sales

```
daily_sales = df.groupby('date')['item_count'].sum().reset_index()
```

```
daily_sales['year'] = daily_sales['date'].dt.year
```

```
daily_sales['month'] = daily_sales['date'].dt.month
```

```
daily_sales['day'] = daily_sales['date'].dt.day
```

```
daily_sales['weekday'] = daily_sales['date'].dt.dayofweek
```

```
daily_sales['quarter'] = daily_sales['date'].dt.quarter
```

Train/Test Split: Last 6 months as test

```
cutoff = daily_sales['date'].max() - pd.DateOffset(months=6)
```

```
train = daily_sales[daily_sales['date'] <= cutoff]
```

```
test = daily_sales[daily_sales['date'] > cutoff]
```

```
features = ['year', 'month', 'day', 'weekday', 'quarter']
```

```
X_train = train[features]
```

```
y_train = train['item_count']
```

```
X_test = test[features]
```

```
y_test = test['item_count']
```

Models

```
models = {
```

```
    "Linear Regression": LinearRegression(),
```

```
    "Random Forest": RandomForestRegressor(n_estimators=100, random_state=42),
```

```
    "XGBoost": XGBRegressor(n_estimators=100, random_state=42)
```

```
}
```

```

print("\nModel Evaluation (RMSE):")
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    rmse = mean_squared_error(y_test, preds, squared=False)
    print(f"{name}: RMSE = {rmse:.2f}")

# Use best model (assume Random Forest)
best_model = RandomForestRegressor(n_estimators=100, random_state=42)
best_model.fit(X_train, y_train)

# Forecast next 365 days
future_dates = pd.date_range(start=daily_sales['date'].max() + pd.Timedelta(days=1),
periods=365)
future_df = pd.DataFrame({'date': future_dates})
future_df['year'] = future_df['date'].dt.year
future_df['month'] = future_df['date'].dt.month
future_df['day'] = future_df['date'].dt.day
future_df['weekday'] = future_df['date'].dt.dayofweek
future_df['quarter'] = future_df['date'].dt.quarter

# Predict
future_df['predicted_sales'] = best_model.predict(future_df[features])

# Plot historical + future forecast
plt.figure(figsize=(14,6))
plt.plot(daily_sales['date'], daily_sales['item_count'], label="Historical Sales")
plt.plot(future_df['date'], future_df['predicted_sales'], label="Forecasted Sales", linestyle='--')

```

```
plt.title("Sales Forecast for Next Year")
```

```
plt.xlabel("Date")
```

```
plt.ylabel("Items Sold")
```

```
plt.legend()
```

```
plt.tight_layout()
```

```
plt.show()
```

Screenshots:

```
[38] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor

from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
```

```
[39] df_rest = pd.read_csv("restaurants.csv")
df_items = pd.read_csv("items.csv")
df_sales = pd.read_csv("sales.csv")
```

```
[40] df_items.rename(columns={'name': 'item_name'}, inplace=True)
df_rest.rename(columns={'name': 'store_name'}, inplace=True)
# df_sales.rename(columns={'name': 'item_name'}, inplace=True) # This line caused the error as df_sales does not have a 'name' column
```

```
[41] print("Sales Columns:", df_sales.columns)
print("Items Columns:", df_items.columns)
print("Restaurants Columns:", df_rest.columns)
```


```
Sales Columns: Index(['date', 'item_id', 'price', 'item_count'], dtype='object')
Items Columns: Index(['id', 'store_id', 'item_name', 'kcal', 'cost'], dtype='object')
Restaurants Columns: Index(['id', 'store_name'], dtype='object')
```

```
[42] df = df_sales.merge(df_items, left_on='item_id', right_on='id', how='left') \
      .merge(df_rest, left_on="store_id", right_on="id", suffixes=('_item', '_store'))
```


```
[43] df['date'] = pd.to_datetime(df['date'])
df.dropna(inplace=True)
```

```
[44] df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
df['day'] = df['date'].dt.day
df['weekday'] = df['date'].dt.dayofweek
df['quarter'] = df['date'].dt.quarter
```

```
[45] plt.figure(figsize=(12,5))
df.groupby('date')['item_count'].sum().plot(title="Daily Sales Trend")
plt.xlabel("Date")
plt.ylabel("Total Items Sold")
plt.tight_layout()
plt.show()
```



```
[46] plt.figure(figsize=(8,4))
sns.barplot(data=df, x='weekday', y='item_count', estimator=sum)
plt.title("Sales by Weekday")
plt.xlabel("Day of Week (0=Monday)")
plt.ylabel("Total Items Sold")
plt.tight_layout()
plt.show()
```



```
[47] top_items = df.groupby('item_name')['item_count'].sum().sort_values(ascending=False).head(10)
print("Top 10 Items Sold:\n", top_items)
```

```
Top 10 Items Sold:
item_name
Strawberry Smoothy      236337.0
Frozen Milky Smoothy    103263.0
Amazing pork lunch      61043.0
Mutton Dinner           52772.0
Orange Juice            43874.0
Blue Ribbon Beef Entree 42774.0
Amazing Steak Dinner with Rolls 34439.0
Sweet Frozen Soft Drink 27490.0
Sea Bass with Vegetables Dinner 23839.0
Sweet Lamb Cake         18764.0
Name: item_count, dtype: float64
```

```
# Most profitable store
df['revenue'] = df['item_count'] * df['price']
top_store = df.groupby('store_name')['revenue'].sum().sort_values(ascending=False)
print("Most Profitable Stores:\n", top_store)
```

```
Most Profitable Stores:
store_name
Bob's Diner      6337275.69
Fou Cher         27885.37
Corner Cafe      16551.43
Surfs Up         15651.49
Beachfront Bar   3796.20
Sweet Shack      2578.27
Name: revenue, dtype: float64
```



```
[49] # Most expensive item per store
expensive_items = df.groupby('store_name').apply(lambda x: x.sort_values('cost', ascending=False).iloc[0])
print("Most Expensive Item per Store:\n", expensive_items[['item_name', 'cost', 'kcal']])
```

Most Expensive Item per Store:

store_name	item_name	cost	kcal
Beachfront Bar	Sweet Vegi Soft Drink	5.70	538
Bob's Diner	Sweet Fruity Cake	29.22	931
Corner Cafe	Pike Lunch	26.37	653
Fou Cher	Blue Ribbon Fruity Vegi Lunch	53.98	881
Surfs Up	Steak Meal	26.21	607
Sweet Shack	Blue Ribbon Frozen Milky Cake	7.70	636

/tmp/ipython-input-49-3444379855.py:2: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, expensive_items = df.groupby('store_name').apply(lambda x: x.sort_values('cost', ascending=False).iloc[0])

```
[50] # 8. Modeling and Forecasting

# Aggregate daily sales
daily_sales = df.groupby('date')['item_count'].sum().reset_index()
daily_sales['year'] = daily_sales['date'].dt.year
daily_sales['month'] = daily_sales['date'].dt.month
daily_sales['day'] = daily_sales['date'].dt.day
daily_sales['weekday'] = daily_sales['date'].dt.dayofweek
daily_sales['quarter'] = daily_sales['date'].dt.quarter

# Train/Test Split: Last 6 months as test
cutoff = daily_sales['date'].max() - pd.DateOffset(months=6)
train = daily_sales[daily_sales['date'] <= cutoff]
test = daily_sales[daily_sales['date'] > cutoff]

features = ['year', 'month', 'day', 'weekday', 'quarter']
X_train = train[features]
y_train = train['item_count']
X_test = test[features]
y_test = test['item_count']
```

```
[52] # Models
models = {
    "Linear Regression": LinearRegression(),
    "Random Forest": RandomForestRegressor(n_estimators=100, random_state=42),
    "XGBoost": XGBRegressor(n_estimators=100, random_state=42)
}

print("\nModel Evaluation (RMSE):")
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
    print(f"{name}: RMSE = {rmse:.2f}")
```

Model Evaluation (RMSE):

Linear Regression: RMSE = 246.54

Random Forest: RMSE = 62.16

XGBoost: RMSE = 61.24

```
[53] # Use best model (assume Random Forest)
best_model = RandomForestRegressor(n_estimators=100, random_state=42)
best_model.fit(X_train, y_train)
```

```
RandomForestRegressor
RandomForestRegressor(random_state=42)
```

```
[54] # Forecast next 365 days
future_dates = pd.date_range(start=daily_sales['date'].max() + pd.Timedelta(days=1), periods=365)
future_df = pd.DataFrame({'date': future_dates})
future_df['year'] = future_df['date'].dt.year
future_df['month'] = future_df['date'].dt.month
future_df['day'] = future_df['date'].dt.day
future_df['weekday'] = future_df['date'].dt.dayofweek
future_df['quarter'] = future_df['date'].dt.quarter
```

```
[55] # Predict
future_df['predicted_sales'] = best_model.predict(future_df[features])
```

```
[56] # Plot historical + future forecast
plt.figure(figsize=(14,6))
plt.plot(daily_sales['date'], daily_sales['item_count'], label="Historical Sales")
plt.plot(future_df['date'], future_df['predicted_sales'], label="Forecasted Sales", linestyle='--')
plt.title("Sales Forecast for Next Year")
plt.xlabel("Date")
plt.ylabel("Items Sold")
plt.legend()
plt.tight_layout()
plt.show()
```

